



Statistical Learning 1 (ECE 271A)
HW2

- (a) Using the training data in TrainingSamplesDCT 8.mat compute the histogram estimate of the prior $P_Y(i), i \in \{cheetah, grass\}$. Using the results of problem 2 compute the maximum likelihood estimate for the prior probabilities. Compare the result with the estimates that you obtained last week. If they are the same, interpret what you did last week. If they are different, explain the differences.

The maximum likelihood estimate for the prior probabilities is the same as the previous estimation last week.

size of total training sample = size of cheetah training sample + size of grass training sample

$$P_Y(cheetah) = \frac{\text{size of cheetah training sample}}{\text{size of total training sample}}$$

$$P_Y(grass) = \frac{\text{size of grass training sample}}{\text{size of total training sample}}$$

In this case, $P_Y(cheetah) = 0.1919, P_Y(grass) = 0.8081$.

- (b) Using the training data in TrainingSamplesDCT 8.mat, compute the maximum likelihood estimates for the parameters of the class conditional densities $P_{(X|Y)}(x|cheetah)$ and $P_{(X|Y)}(x|grass)$ under the Gaussian assumption. Denoting by $X = \{X_1, \dots, X_{64}\}$ the vector of DCT coefficients, create 64 plots with the marginal densities for the two classes - $P_{(X_k|Y)}(x_k|cheetah)$ and $P_{(X_k|Y)}(x_k|grass)$, $k = 1, \dots, 64$ - on each. Use different line styles for each marginal. Select, by visual inspection, what you think are the best 8 features for classification purposes and what you think are the worst 8 features (you can use the subplot command to compare several plots at a time). Hand in the plots of the marginal densities for the best-8 and worst-8 features (once again you can use subplot, this should not require more than two sheets of paper). In each subplot indicate the feature that it refers to.


Given the marginal densities of the two classes - $P_{(X_k|Y)}(x_k|cheetah)$ and $P_{(X_k|Y)}(x_k|grass)$ are Gaussian assumption, we can plot the densities and compute the maximum likelihood by the following equation.

estimated mean = sample mean

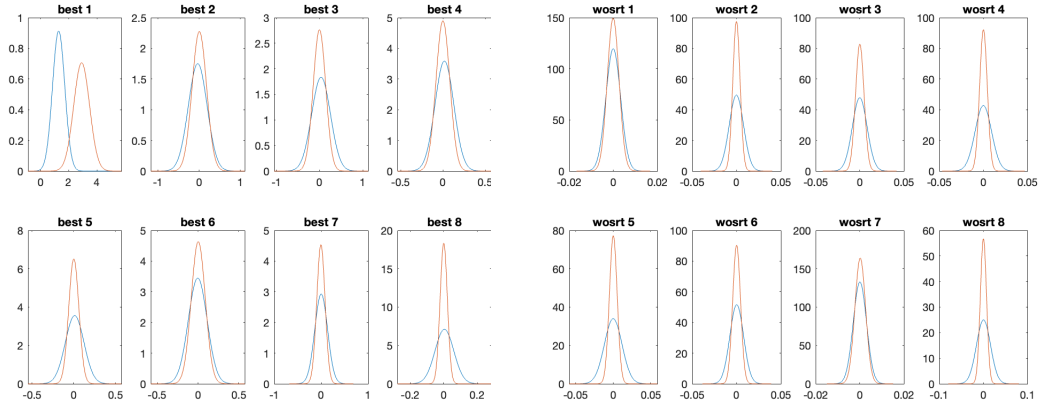
estimated std = sample std

$$L = \frac{1}{(\sigma_T \sqrt{2\pi})^N} e^{-\frac{1}{2} \sum_{i=1}^N \left(\frac{T_i - \bar{T}}{\sigma_T}\right)^2} \quad \text{In this case, } N = 64.$$

The following graphs are the densities of all the features, the best 8 features, and the worst 8 features. We can notice that the better feature is, the more distance between the mean of the foreground and the background is.



(a) All features



(b) Best 8 features

(c) Worst 8 features

- (c) Compute the Bayesian decision rule and classify the locations of the cheetah image using i) the 64-dimensional Gaussians, and ii) the 8-dimensional Gaussians associated with the best 8 features. For the two cases, plot the classification masks and compute the probability of error by comparing with cheetah mask.bmp. Can you explain the results?


For this problem, we could use the DCT process as HW1. While determining the classification, we use the Gaussian classifier. The following 5 steps are my approach.

- 1) To normalize the value in the range $[0, 1]$, divide `img cheetah.bmp` by 255.
- 2) Padding the `img cheetah.bmp` with 0 for all four sides to get a new image size with (263, 278). After performing sliding windows, we could get the matrix with the size of (255, 270), the same as the input `img`, by padding.
- 3) For each 8×8 block, we compute DCT, order coefficients with zig-zag scan.
- 4) Determine the predicted class in the current block by Gaussian classifier. The BDR is


$$i^*(x) = \underset{i}{\operatorname{argmax}} \left[-\frac{1}{2}(x - \mu_i)^T \sum_i^{-1} (x - \mu_i) - \frac{1}{2} \log(2\pi)^d |\sum_i| + \log P_Y(i) \right]$$

where x denotes the flatten vector from each DCT sliding blocks, while we can get \sum_i and μ_i from calculation. The predicted class = 0 (background) if the value with $i = 0$ in the argmax function is larger than that with $i = 1$ in the argmax function. Likewise, The predicted class = 1 (foreground) if the value with $i = 1$ in the argmax function is larger than or equal to that with $i = 0$ in the argmax function.

- 5) Create a binary mask *all* with 1's for cheetah (foreground) and 0's for grass (background).
- 6) Repeat step 1. to step 5. with the best 8 features and create a binary mask *best8*.



(a) All features



(b) Best 8 features

For all the pixel in the binary mask *all* and *best8*, we could check whether the predicted label equals to the ground truth label in *cheetah_mask.bmp*. The error rate could be calculated by

$$\text{error rate} = \frac{\text{number of different pixel value}}{\text{number of pixels in the mask}}$$

In this case, error rate with all features = **0.1357** and error rate with best 8 features = **0.1000**. Both the values are smaller than that we obtained from HW1.

Code

```
HW2_1.m HW2_2.m HW2_3.m HW1_4.m +
1 load('TrainingSamplesDCT_8_new.mat');
2 total = size(TrainsampleDCT_FG) + size(TrainsampleDCT_BG);
3 priorFG = length(TrainsampleDCT_FG) / total(1);
4 priorBG = length(TrainsampleDCT_BG) / total(1);
```

Figure 3: Code for problem (a)

```

HW2_1.m HW2_2.m HW2_3.m HW1_4.m +
1 % Compute the MLE for the foreground;
2 fSize = size(TrainsampleDCT_FG);
3 fSize = fSize(1);
4 fBase = zeros(1, 64);
5 fTmp = zeros(1, 64);
6 fMean = mean(TrainsampleDCT_FG, 1);
7 fVar = var(TrainsampleDCT_FG, 1);
8 fStd = std(TrainsampleDCT_FG, 1);
9
10 for col = 1 : 64
11     fBase(col) = (-2/fSize) * log(2 * pi);
12 end
13
14 for col = 1 : 64
15     for row = 1 : fSize
16         fTmp(col) = fTmp(col) + ((TrainsampleDCT_FG(row, col) - fMean(col)) / fStd(col))^2;
17     end
18     fTmp(col) = 0.5 * fTmp(col);
19 end
20
21 fMLELog = fBase - fSize * log(fStd) - fTmp;
22
23
24 % Compute the MLE for the background;
25 bSize = size(TrainsampleDCT_BG);
26 bSize = bSize(1);
27 bBase = zeros(1, 64);
28 bTmp = zeros(1, 64);
29 bMean = mean(TrainsampleDCT_BG, 1);
30 bVar = var(TrainsampleDCT_BG, 1);
31 bStd = std(TrainsampleDCT_BG, 1);
32
33 for col = 1 : 64
34     bBase(col) = (-2/bSize) * log(2 * pi);
35 end
36
37 for col = 1 : 64
38     for row = 1 : bSize
39         bTmp(col) = bTmp(col) + ((TrainsampleDCT_BG(row, col) - bMean(col)) / bStd(col))^2;
40     end
41     bTmp(col) = 0.5 * bTmp(col);
42 end
43
44 bMLELog = bBase - bSize * log(bStd) - bTmp;
45
46 % plot the images;
47 sampleNum = 100;
48 stdNum = 5;
49
50 distance = zeros(1, 64);
51
52 for i = 1 : 64
53     subplot(8, 8, i);
54     fX = linspace(fMean(i) - stdNum * fStd(i), fMean(i) + stdNum * fStd(i), sampleNum);
55     bX = linspace(bMean(i) - stdNum * bStd(i), bMean(i) + stdNum * bStd(i), sampleNum);
56     x = sort([fX, bX]);
57
58     fY = normpdf(x, fMean(i), fStd(i));
59     bY = normpdf(x, bMean(i), bStd(i));
60     distance(i) = abs(fMean(i) - bMean(i));
61
62     plot(x, fY, x, bY);
63 end
64 savefig('all_features.fig');
65
66 % plot the best and the worst feature;
67 [out, idx] = sort(distance, 'descend');
68 best = idx(1 : 8);
69 worst = flipr(idx(57 : 64));
70
71 % best;
72 for i = 1 : size(best, 2)
73     subplot(2, 4, i);
74     fX = linspace(fMean(best(i)) - stdNum * fStd(best(i)), fMean(best(i)) + stdNum * fStd(best(i)), sampleNum);
75     bX = linspace(bMean(best(i)) - stdNum * bStd(best(i)), bMean(best(i)) + stdNum * bStd(best(i)), sampleNum);
76     x = sort([fX, bX]);
77
78     fY = normpdf(x, fMean(best(i)), fStd(best(i)));
79     bY = normpdf(x, bMean(best(i)), bStd(best(i)));
80
81     plot(x, fY, x, bY);
82     title(['best ', num2str(i)]);
83 end
84 savefig('best 8 features.fig');
85
86 % worst;
87 for i = 1 : size(worst, 2)
88     subplot(2, 4, i);
89     fX = linspace(fMean(worst(i)) - stdNum * fStd(worst(i)), fMean(worst(i)) + stdNum * fStd(worst(i)), sampleNum);
90     bX = linspace(bMean(worst(i)) - stdNum * bStd(worst(i)), bMean(worst(i)) + stdNum * bStd(worst(i)), sampleNum);
91     x = sort([fX, bX]);
92
93     fY = normpdf(x, fMean(worst(i)), fStd(worst(i)));
94     bY = normpdf(x, bMean(worst(i)), bStd(worst(i)));
95
96     plot(x, fY, x, bY);
97     title(['worst ', num2str(i)]);
98 end

```

Figure 4: Code for problem (b)

```

HW2_1.m HW2_2.m HW2_3.m HW1_4.m +
1 % read img;
2 img = imread('cheetah.bmp');
3 img = double(img) / 255;
4
5 % padding img with zero;
6 I = zeros(263, 278);
7 for row = 5 : 259
8     for col = 5 : 274
9         I(row, col) = img(row - 4, col - 4);
10    end
11 end
12
13 % read pattern;
14 pattern = readmatrix('Zig-Zag Pattern.txt');
15
16 % create the bitmask by all features;
17 all = zeros(255, 270);
18
19 fCov = cov(TrainSampleDCT_FG);
20 bCov = cov(TrainSampleDCT_BG);
21 fCovInv = inv(fCov);
22 bCovInv = inv(bCov);
23 fCovDet = det(fCov);
24 bCovDet = det(bCov);
25 fAllMean = mean(TrainSampleDCT_FG, 1);
26 bAllMean = mean(TrainSampleDCT_BG, 1);
27
28 for row = 1 : 255
29     for col = 1 : 270
30         block = zeros(8, 8);
31         % get the block;
32         for r = row : row + 7
33             for c = col : col + 7
34                 block(r - row + 1, c - col + 1) = I(r, c);
35             end
36         end
37
38         % compute DCT;
39         dct2Block = dct2(block);
40         flatBlock = zeros(1, 64);
41
42         % order coefficients with zig-zag scan;
43         for r = 1 : 8
44             for c = 1 : 8
45                 flatBlock(1, pattern(r, c) + 1) = dct2Block(r, c);
46             end
47         end
48     end
49 end
50 % use Gaussian classifier to find class Y for each block;
51
52 fRes = -0.5 * (flatBlock - fAllMean) * fCovInv * transpose(flatBlock - fAllMean) - ...
53     0.5 * log(2 * pi) ^ 64 * fCovDet + log(priorFG);
54
55 bRes = -0.5 * (flatBlock - bAllMean) * bCovInv * transpose(flatBlock - bAllMean) - ...
56     0.5 * log(2 * pi) ^ 64 * bCovDet + log(priorBG);
57
58 % create a binary mask;
59 if fRes >= bRes
60     all(row, col) = 1;
61 else
62     all(row, col) = 0;
63 end
64 end
65
66 imshow(uint8(all), [0 1]);
67 savefig('all.fig');
68
69 % create the bitmask by best 8 features;
70 best8 = zeros(255, 270);
71
72 fCov = cov(TrainSampleDCT_FG(:, best));
73 bCov = cov(TrainSampleDCT_BG(:, best));
74 fCovInv = inv(fCov);
75 bCovInv = inv(bCov);
76 fCovDet = det(fCov);
77 bCovDet = det(bCov);
78 fBestMean = mean(TrainSampleDCT_FG(:, best), 1);
79 bBestMean = mean(TrainSampleDCT_BG(:, best), 1);
80
81 for row = 1 : 255
82     for col = 1 : 270
83         block = zeros(8, 8);
84         % get the block;
85         for r = row : row + 7
86             for c = col : col + 7
87                 block(r - row + 1, c - col + 1) = I(r, c);
88             end
89         end
90
91         % compute DCT;
92         dct2Block = dct2(block);
93         flatBlock = zeros(1, 64);
94
95         % order coefficients with zig-zag scan;
96         for r = 1 : 8
97             for c = 1 : 8
98                 flatBlock(1, pattern(r, c) + 1) = dct2Block(r, c);
99             end
100         end
101
102         % use Gaussian classifier to find class Y for each block;
103         fRes = -0.5 * (flatBlock - fBestMean) * fCovInv * transpose(flatBlock - fBestMean) - ...
104             0.5 * log(2 * pi) ^ 64 * fCovDet + log(priorFG);
105
106         bRes = -0.5 * (flatBlock - bBestMean) * bCovInv * transpose(flatBlock - bBestMean) - ...
107             0.5 * log(2 * pi) ^ 64 * bCovDet + log(priorBG);
108
109         % create a binary mask;
110         if fRes >= bRes
111             best8(row, col) = 1;
112         else
113             best8(row, col) = 0;
114         end
115     end
116 end
117
118 imshow(uint8(best8), [0 1]);
119 savefig('best8.fig');
120
121 % compute the error rate;
122 mask = imread('cheetah_mask.bmp');
123 allErrorCount = 0;
124 bestErrorCount = 0;
125 sizes = size(mask);
126 rows = sizes(1);
127 cols = sizes(2);
128
129 % check whether the predicted label equals to the ground truth label;
130 for row = 1 : rows
131     for col = 1 : cols
132         if (mask(row, col) / 255 ~= all(row, col))
133             allErrorCount = allErrorCount + 1;
134         end
135         if (mask(row, col) / 255 ~= best8(row, col))
136             bestErrorCount = bestErrorCount + 1;
137         end
138     end
139 end
140
141 allErrorRate = allErrorCount / (rows * cols);
142 disp(allErrorRate);
143 bestErrorRate = bestErrorCount / (rows * cols);
144 disp(bestErrorRate);

```

Figure 5: Code for problem (c)